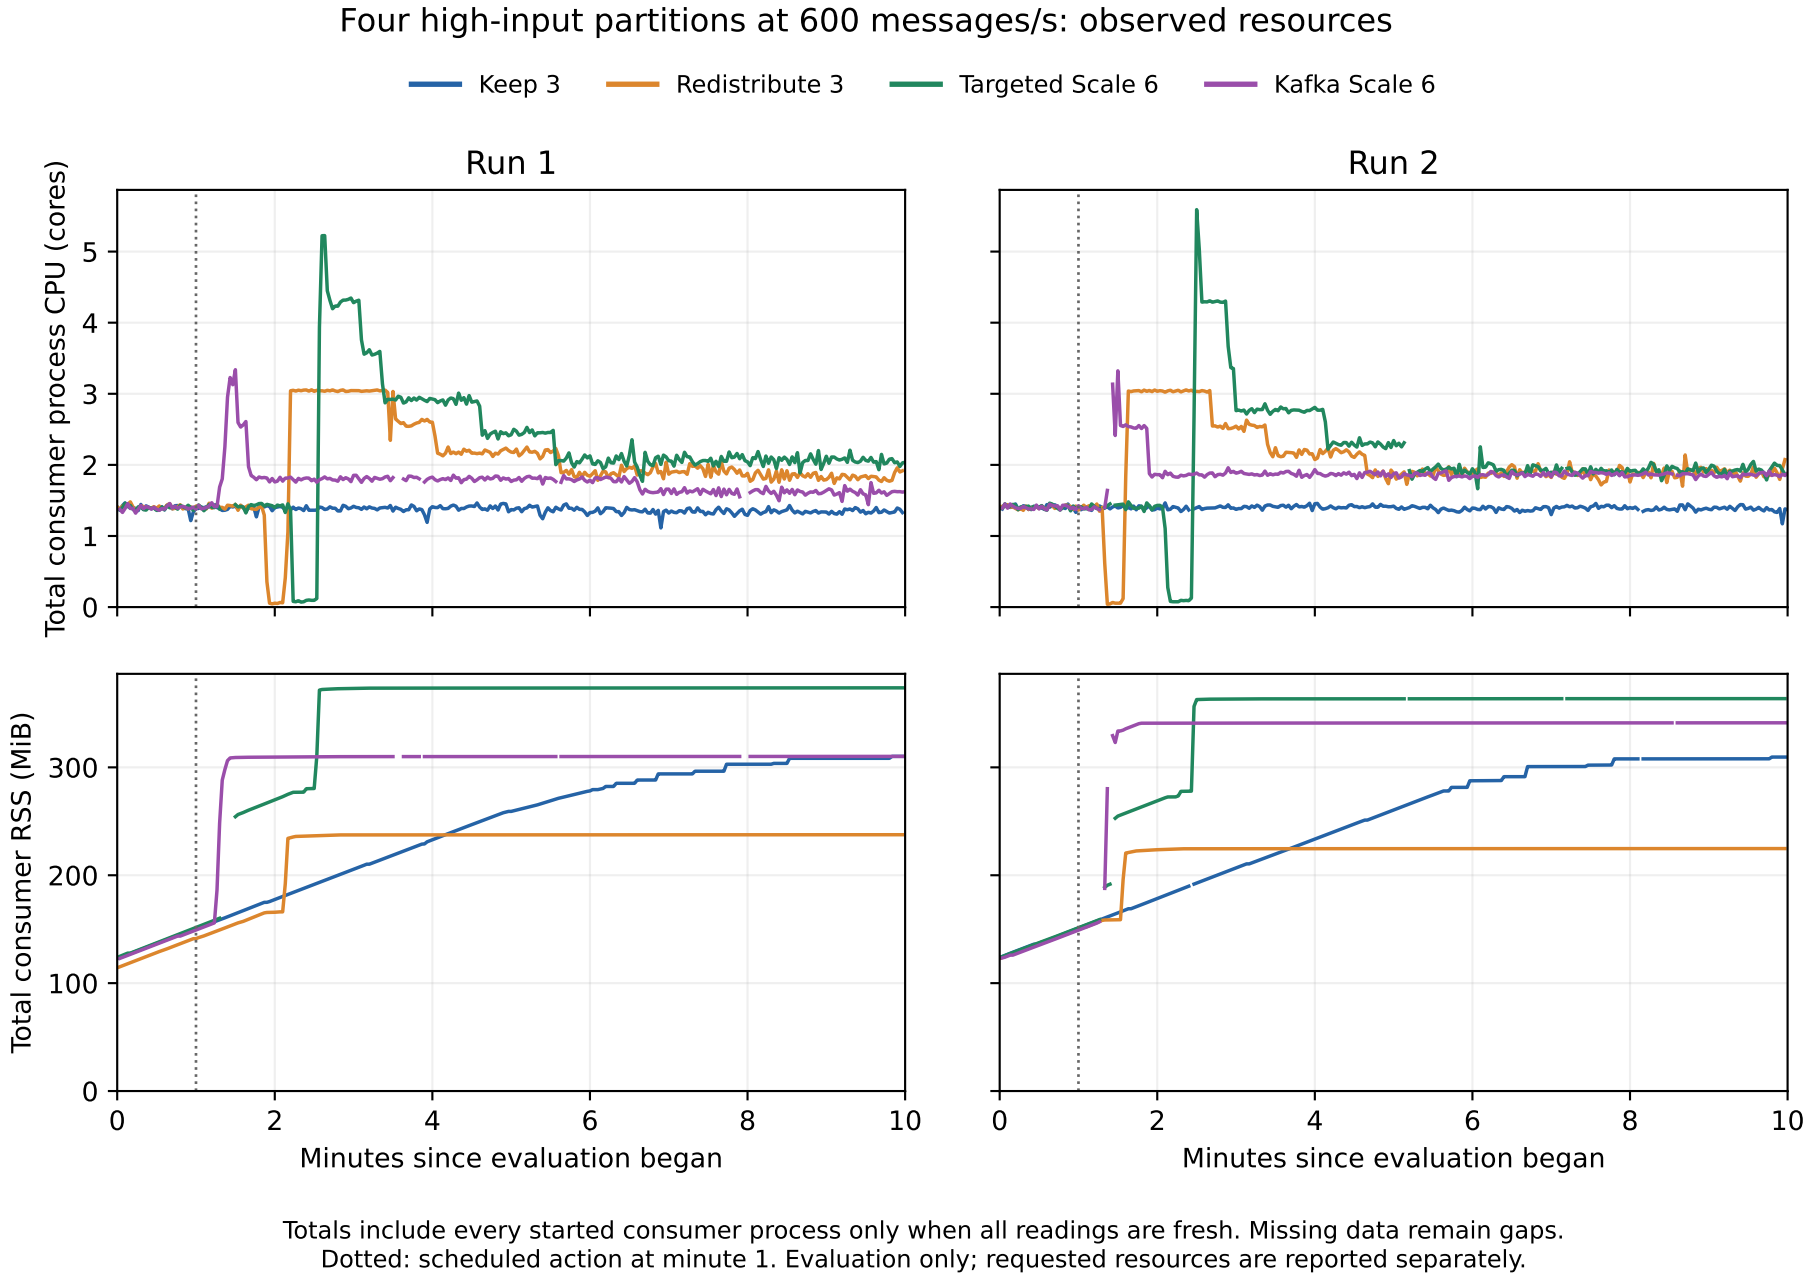
Four high-input partitions at 600 messages/s: observed resources
Keep 3
Redistribute 3
Targeted Scale 6
Kafka Scale 6
Run 1
Run 2
Total consumer process CPU (cores)
Total consumer RSS (MiB)
Minutes since evaluation began
Minutes since evaluation began
Totals include every started consumer process only when all readings are fresh. Missing data remain gaps.
Dotted: scheduled action at minute 1. Evaluation only; requested resources are reported separately.