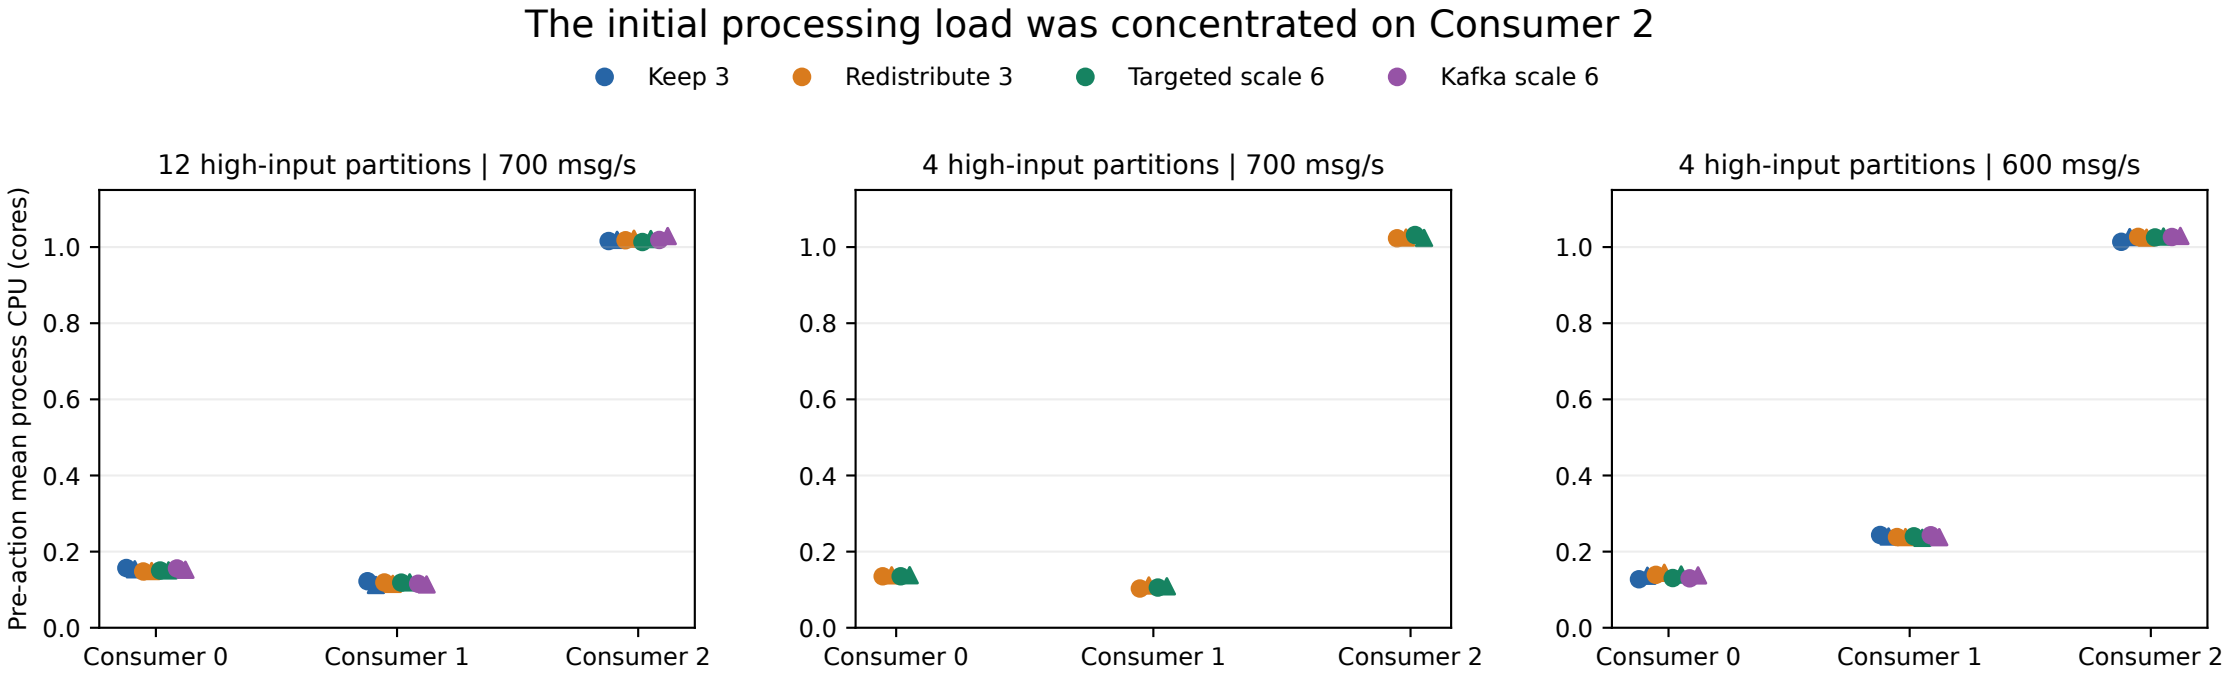
The initial processing load was concentrated on Consumer 2
Keep 3
Redistribute 3
Targeted scale 6
Kafka scale 6
12 high-input partitions | 700 msg/s
4 high-input partitions | 700 msg/s
4 high-input partitions | 600 msg/s
Pre-action mean process CPU (cores)
0.0
0.2
0.4
0.6
0.8
1.0
Consumer 0
Consumer 1
Consumer 2
Arithmetic means at evaluation seconds 30, 32, ..., 58. Small horizontal offsets separate trials visually.
All selected high-input partitions initially belonged to Consumer 2. Low CPU on other consumers suggests uneven use,
but CPU observations do not measure how much additional work those consumers can sustainably process.